

House Price Prediction using Linear Regression

Overview

This project demonstrates how to predict house prices using **Linear Regression**. By leveraging features like the number of bedrooms, square footage, and more, the model aims to make accurate price predictions.

The project walks through data exploration, training a linear regression model, evaluating its performance, and visualizing the results.

Project Structure

house-price-prediction/

|

|— house_prices.csv # The dataset used for model training

|— house_price_prediction.py # Main script for training and evaluating the model

|— requirements.txt # Required packages and libraries

|— README.md # Project documentation

|— LICENSE # License for the project

Dataset

The dataset consists of 4,600 records with the following features:

- date: Date of the sale
- price: Sale price of the house (target variable)
- bedrooms: Number of bedrooms
- bathrooms: Number of bathrooms

- sqft_living: Square footage of living space
- sqft_lot: Square footage of the lot
- floors: Number of floors
- waterfront: Indicates if the property is waterfront (1) or not (0)
- view: Rating of the house's view (0-4)
- condition: Condition of the house (1-5)
- sqft_above: Square footage of the house above the basement
- sqft_basement: Square footage of the basement
- yr_built: Year the house was built
- yr_renovated: Year the house was renovated

Requirements

- Python 3.x
- Pandas
- NumPy
- Scikit-Learn
- Matplotlib
- Seaborn

To install the required dependencies, use:

```
pip install -r requirements.txt
```

Installation

1. Clone the repository:

```
git clone https://github.com/aswinrajkt586/House-Price-Prediction.git
```

2. Navigate to the project directory:

```
cd House-Price-Prediction
```

3. Install the necessary dependencies:

```
pip install -r requirements.txt
```

Usage

To train and evaluate the linear regression model, run the following:

```
house_price_prediction.py
```

The script will:

- Load and preprocess the dataset.
- Train the Linear Regression model.
- Evaluate model performance on test data.
- Output Mean Squared Error (MSE) and (R^2).

Model Evaluation

- Mean Squared Error (MSE): Evaluates the average squared difference between predicted and actual values.
- R^2 score: Represents how much variance in the target variable (house prices) is explained by the model.

Results

The linear regression model achieved:

- MSE: 111408.9
- R^2 score: 47.8

Future Work

Potential improvements:

- Implementing advanced models like Random Forest or Gradient Boosting.
- Feature engineering to improve model accuracy.
- Hyperparameter tuning for better performance.

License

This project is licensed under the MIT License. See the LICENSE file for more details.